

How LLMs Feel About Applicants: Capturing Social Intuition as Vectors

Emrecaan Ulu
University of Konstanz
emrecaan.ulu@uni-konstanz.de

Jorge Lastra Cerda
University of Konstanz
jorge.lastra-cerda@uni-konstanz.de

Abstract

Recent interpretability work shows that emotions can be represented as linear directions in a language model’s residual stream. We adapt this approach to social perception across nine open-weight models from the Gemma, Qwen, and Llama families. Using synthetic stories, we derive warmth and competence directions and then test whether name-level projections align with human ratings and whether steering these directions changes callback recommendations. Human alignment was model-dependent: both dimensions aligned positively in Gemma-3 and Qwen3.6, whereas Llama-3.1 and Qwen3-14B inverted warmth, and Gemma-4 showed mixed alignment across checkpoints. Hiring steering was similarly heterogeneous: positive interventions increased callback margins in some checkpoints, decreased them in others, and produced nonlinear or decision-stable responses elsewhere. Thus, a common extraction procedure recovers warmth and competence directions across models, but their correspondence with human perceptions and causal influence on hiring decisions are not uniform.

Introduction

Work in progress.

Background: Reading a Concept Out of a Language Model

From Layers to Directions. A large language model processes a piece of text by passing it through a long stack of layers, and each layer reads from and writes back to a shared running representation known as the residual stream (Vaswani et al., 2017; Elhage et al., 2021). Think of the residual stream as a notebook passed from one layer to the next: every layer reads what earlier layers have written, adds its own notes, and hands the notebook onward. At any chosen depth, this running state is simply a long list of numbers, a high-dimensional vector, and circuit-level analyses have shown that specific features of the input, from low-level syntax to high-level semantic content, leave consistent geometric traces in that vector across many different contexts (Olah et al., 2020). The idea that such traces take the form of *directions* rather than isolated coordinates goes back to word-embedding analogies, where arithmetic on word vectors captures semantic relationships (Mikolov et al., 2013), and the *linear representation hypothesis* extends this claim to modern transformer models: a high-level feature corresponds to a direction in the residual stream, and projecting the model’s state onto that direction yields a scalar measuring how strongly the feature is present (Park et al., 2024). Extracting a direction for some target concept

is then a matter of contrast: we collect residual-stream states from texts that express the concept (condition A) and from matched texts that do not (condition B), and take the mean difference between the two groups,

$$v = \bar{h}_A - \bar{h}_B. \quad (1)$$

Normalizing v to a unit vector $\hat{v}$ gives a concept direction, and the concept score of any new piece of text is the projection $\hat{v}^\top h$. The claim worth testing, and the one we validate here through cross-validated classification and topic-holdout accuracy, is that a direction built this way generalizes beyond the handful of texts used to construct it.

Emotion Vectors. A direction built purely from the mean-difference of Eq. 1 only tells us that high- and low-condition texts sit in different places; it does not yet tell us whether the model actually relies on that direction when it acts. Testing that requires an intervention rather than an observation: at a chosen layer, we can override the model’s own internal state during inference,

$$h' = h + \alpha \hat{v}, \quad (2)$$

where α is a signed scalar that sets how strongly, and in which direction, we push. If the model’s downstream behavior then shifts systematically as α changes, while an unrelated, orthogonal direction produces no comparable shift, the direction is not merely correlated with the concept but participates causally in the model’s computation. Activation addition (Turner et al., 2023) and representation engineering (Zou et al., 2023) established this intervene-and-observe recipe and




Figure 1: Emotion vectors as directions in activation space (schematic). Conceptually, an emotion vector starts from a neutral model state h . A prompt that evokes a given emotion, such as fear, pushes the residual stream toward a direction $\hat{v}_{\text{fear}}$, encoded here with its representative d -dimensional coordinates. This direction behaves like a commuter who takes the same route to work every day: the attitude, direction, and flow stay the same each time. Across thousands of different prompts that express the same emotion in different words, the model converges on this same direction through its layers, and it produces its output by moving along that direction. After Sofroniew et al. (2026), who call this an emotion vector; we reuse the geometry, not the construct.

demonstrated it across a range of models and behaviors. Most recently, Sofroniew et al. (2026) applied exactly this recipe to emotion concepts in Claude Sonnet: building contrastive text pairs for each emotion, they extracted mean-difference directions in the residual stream, pushed the model’s state along those directions during inference, and observed systematic, interpretable shifts in the model’s behavior. Their study also introduced a neutral-corpus PCA denoising procedure that strips out general evaluative valence from a direction vector, a step we adopt to separate warmth and competence content from shared narrative polarity (detailed in the Supplementary Materials). Fig. 1 illustrates the resulting picture: each emotion occupies its own direction in residual-stream space, with a numeric representation that can be extracted, inspected, and used to steer the model.

From Emotions to Hiring. Nothing in this read-and-steer recipe is specific to emotion. Any construct that is linearly encoded in the residual stream should, in principle, be extractable by the same mean-difference procedure and testable by the same steering intervention, whether that construct is an emotional state or a dimension of social perception. Hiring is a domain where such a construct would matter immediately: be-

havioral audits have already documented that large language models produce uneven callback recommendations across demographic groups (An et al., 2024; Gaebler et al., 2024), yet none of that work locates the representation inside the model that might be doing the work. We therefore ask whether a social-perception construct known to predict human callback outcomes is encoded the same way emotion concepts are, and whether steering it changes a model’s own callback recommendations. The next section describes how we chose that construct and built the materials needed to test it.

Adapting the Method to Hiring

Why Warmth and Competence. The Stereotype Content Model proposes that social perception organizes along two fundamental axes: warmth, capturing perceived intentions and trustworthiness, and competence, capturing perceived ability and effectiveness (Fiske et al., 2002). Evidence across cultural contexts suggests that warmth and competence are recurring dimensions of both interpersonal impressions and perceptions of social groups (Fiske et al., 2007; Cuddy et al., 2008). Warmth and competence therefore give us two well-established, independently motivated axes to search for inside a language model, rather than an ad hoc pair invented for this study.

The Link to Hiring. Gallo et al. (2024) provide the empirical bridge that connects these two axes to real labor-market outcomes. Their released ratings data contain 24,220 warmth and competence judgments from 787 raters covering 282 applicant first names drawn from ten source studies. For the eight name-based correspondence studies with usable callback data, a more favorable first principal component combining warmth and competence was moderately associated with higher callback rates on average. The effect nevertheless varied substantially across studies, and its prediction interval included negative relationships. These findings provide a direct, quantitative connection between the social-perception dimensions we probe and measured labor-market outcomes, while showing that the relationship is not uniform. Behavioral audits of LLMs have separately documented name-driven callback disparities (An et al., 2024; An et al., 2025; Gaebler et al., 2024). To our knowledge, these studies have not tested whether internal warmth and competence representations carry that signal. We use the dataset released by Gallo et al. in two ways: as an external validity check on our probes, and as the reference point against which model-level disparities are compared.

Concept Stories Without Identity. Probing warmth and competence directly from hiring prompts

would confound the social-perception dimensions with the hiring decision itself, so the direction vectors have to come from a separate set of non-hiring materials. Identity control is essential here. In unconstrained pilot generations, the language model systematically assigned male, Anglo-named protagonists to low-warmth and low-competence stories, which would have introduced demographic expectations before the hiring evaluation. We therefore use nameless, gender-neutral protagonists so that the vectors reflect depicted behavior rather than the demographic signal tested later.

Applicant Names. The hiring evaluation itself draws on a separate stimulus set: the 282 applicant first names from Gallo et al. (2024), each inserted into an otherwise identical, fixed job application so that the name is the only signal that varies across trials. Of these 282 names, 149 also appear in the published per-name callback dataset from the same source and carry US race and gender labels; this subset forms the benchmark comparison for demographic disparity. The full set of 282 is used for the probe-versus-human alignment check, since every name carries crowdsourced warmth and competence ratings even when published callback data are unavailable.

Methods

Story Preparation. The corpus contains 200 narratives of approximately 90–150 words, divided evenly among high warmth, low warmth, high competence, and low competence (50 stories per condition). Its 50 everyday topics span ten domains, including workplace, learning, community, sport, and travel; each topic appears once in every condition to balance situational vocabulary across poles (Tab. S.6 lists every topic with its word-count range). Following the identity controls described above, every protagonist is nameless and referred to only as “they.” A large language model generated the stories, which were then checked automatically for forbidden vocabulary, length balance, and demographic neutrality. Additional Results, Story Corpus by Topic, gives more detail about the stories: the full topic list, an assigned domain category for each, and per-topic word-count ranges.

Building the Warmth and Competence Vectors. We extracted residual-stream activations from nine open-weights instruction-tuned models spanning three model families and two version generations each for Gemma and Qwen: Gemma-3-12B-it and Gemma-3-27B-it (Gemma Team, 2025), Gemma-4-12B, Gemma-4-26B-A4B, and Gemma-4-31B; Qwen3-14B (Yang et al., 2025), Qwen3.6-27B, and Qwen3.6-35B-A3B; and Llama-3.1-8B-Instruct (Grattafiori et al., 2024). Acti-

vations were read with TransformerLens (Nanda and Bloom, 2022) for the Gemma-3, Qwen3-14B, and Llama-3.1 checkpoints, with its Bridge interface for the three Gemma-4 checkpoints, and with a native Hugging Face backend for the two Qwen3.6 checkpoints, which the released TransformerLens build does not yet support. The *Backend* column of Tab. 1 lists the tool used for each checkpoint. Sofroniew et al. (2026), whose emotion-vector method motivates our approach, describe a progression from token-level emotional connotations in early layers to context-integrated representations in middle-to-late layers. They report most analyses at a layer approximately two thirds through the model, where their evidence indicates that emotion representations encode abstract content relevant to upcoming token predictions. Guided by this result, we operationalized the same relative depth as `probe_layer_frac = 0.66` and fixed it across models before conducting the layer sweeps, rather than selecting a model-specific optimum. Integer rounding placed the probe between 63 and 66 percent of model depth; Tab. 1 reports the exact layer index, model width, and mean residual-stream norm for each model. For each story, activations were mean-pooled across token positions from position 50 to the end of the story, giving one d_{model} -dimensional vector per story. This offset was fixed in advance and applied uniformly across models rather than tuned per architecture: the earliest token positions carry input-agnostic outlier activations that would otherwise dominate the pooled mean (Sun et al., 2024; Xiao et al., 2024), and because the concept stories ran roughly 100 to 170 tokens, starting at position 50 discards about the first third of each narrative so that the pooled state reflects text where the warmth or competence content is already established. High- and low-condition stories cluster in distinct regions of this space (Fig. 2). The warmth and competence direction vectors are simply the mean difference between the two clusters:

$$\begin{aligned} v_W &= \bar{x}_{\text{high warmth}} - \bar{x}_{\text{low warmth}}, \\ v_C &= \bar{x}_{\text{high comp.}} - \bar{x}_{\text{low comp.}} \end{aligned}$$

where $\bar{x}_{\text{high warmth}}$ is the average pooled activation across the 50 high-warmth stories, $\bar{x}_{\text{low warmth}}$ is the same average across the 50 low-warmth stories, and v_W (respectively v_C for competence) is the resulting direction vector. All nine models used the same stories, pooling rule, and prespecified seed (20260527), enabling comparison across architectures.

We check that each direction is a genuine, generalizable signal rather than an artifact of the specific stories used to build it, using five complementary checks:

- **5-fold cross-validated accuracy:** We need to know whether the direction we found is picking up on something real, or whether it just memorized the exact 200 stories we happened to write. The 100 warmth stories

Model	d_{model}	Probe layer	$\ \tilde{h}\ $	Backend
Gemma-3-12B-it	3840	31/48	79,722	TL
Gemma-3-27B-it	5376	40/62	61,576	TL
Gemma-4-12B	3840	31/48	97.8	Bridge
Gemma-4-26B-A4B	2816	19/30	68.4	Bridge
Gemma-4-31B	5376	39/60	188.5	Bridge
Qwen3-14B	5120	26/40	206.6	TL
Qwen3.6-27B	5120	42/64	71.9	HF
Qwen3.6-35B-A3B	2048	26/40	3.6	HF
Llama-3.1-8B-Instruct	4096	20/32	11.4	TL

Table 1: Model architecture and activation scale for nine checkpoints. d_{model} is the number of activation values carried at each token, or the representation width. Probe layer gives the zero-indexed block used to build the directions and the total blocks; 31/48 means index 31 of a 48-block model. $\|\tilde{h}\|$ is the typical activation size there: for each of 200 stories, token activations are averaged and the resulting vector’s length is measured; the table reports the mean of these lengths. This value calibrates steering to each model’s own scale, so larger values do not imply better or stronger models. Backend is the tool used to read activations: TL is TransformerLens, Bridge is its Hugging Face Bridge interface, and HF a native Hugging Face backend used for the Qwen3.6 checkpoints that TransformerLens does not yet support.

(50 high, 50 low) are split into 5 equal groups of 20. In each round, a classifier is trained on 4 of the groups and tested on the 5th group, which it has never seen; this repeats 5 times so every group is tested once. Guessing randomly would get the label right about 50% of the time, so a classifier that gets far more than half of these unseen stories right is reading a real high-versus-low pattern in the stories, not reciting stories it has already seen.

- **Topic-holdout accuracy:** The check above still lets a classifier study some stories about a topic, say a stressful team meeting, before being tested on other stories about that same topic; getting those right could mean it recognized the topic, not the tone. This check closes that gap. The 50 topics are grouped so that every story about a topic, high- and low-condition alike, is held out together, and each test round contains topics the classifier has never read a single sentence about, in either condition. A classifier that still gets these completely unfamiliar topics right is recognizing warmth or competence itself, not the leftover vocabulary of a topic it has already studied.
- **Cohen’s d :** The two checks above only tell us whether high- and low-condition stories can be told apart, not how far apart they really are. Each story is projected onto the direction, reducing it to a single number, and Cohen’s d measures the gap between the high- and low-condition averages on that number line. To know whether that gap is meaningful, we generate 1,000 random directions instead of the real one and compute the same gap for each. Every one of those 1,000 random directions produces a much smaller gap

than our real direction does, which means the separation is not a coincidence.

- **Split-half cosine stability:** A direction built from one arbitrary set of stories should not depend on exactly which stories happened to be written. Each condition’s 50 stories are randomly split into two non-overlapping halves of 25, and a separate direction is built from each half using only that half’s stories. The cosine between these two independently-built directions tells us how similar they are. Across all nine models, this cosine is consistently high (roughly 0.69 to 0.82 for warmth, 0.83 to 0.90 for competence), meaning the direction reflects a stable property of the concept rather than an artifact of which 25 stories happened to end up in each half.
- **Cross-axis classification:** If warmth and competence were fully independent, knowing a story’s position along the warmth direction would tell us nothing about whether it is high- or low-competence. To test this, competence stories are projected onto the warmth direction, and that single number is used to guess the competence label; the same test is then run in the other direction. Across all nine models this guess is highly accurate, typically 0.82 to 1.00, which means the two directions are not fully independent: they share a substantial amount of evaluative content rather than capturing two fully separate constructs.

Two further analyses look beyond a single model’s own directions:

- **Cross-model Spearman agreement:** Nine models, trained independently by different organizations, could each encode warmth as their own unrelated pattern, with no reason for their story rankings to agree with one another. To test this, every pair of the nine models is compared: each story is projected onto both models’ own directions to give one score per model, and the two resulting sets of scores are correlated with Spearman’s rank correlation. All thirty-six pairs show positive agreement on both axes, which means the nine architectures are converging on a shared cross-architecture construct rather than nine independent, unrelated ones.
- **Neutral-corpus PCA denoising:** Some of what warmth and competence have in common in the model has nothing to do with either concept specifically; it looks more like a general tone the model applies regardless of topic. To find and remove that general tone, we ran 1,500 real Wikipedia articles on ordinary, unrelated topics, text with nothing to do with warmth or competence, through the same model and looked at what the warmth and competence directions have in common with those neutral texts. That shared, non-specific part is then subtracted out of both direction vectors, leaving only what is actually specific to warmth and to competence. The Supplementary Ma-

terials (PCA Denoising) describe the full procedure as a robustness check.

Steering the Concept Vectors. Finding a direction that separates high- and low-condition stories does not yet prove the model actually uses that direction when forming its judgment. To test that, we push the model’s internal activity along the direction while it answers, using activation steering (Turner et al., 2023) via TransformerLens hooks at the probe layer (Fig. 2), and see whether its answer changes.

Because different models’ internal activations are different sizes (see the norm column in Tab. 1), we measure how hard we push relative to each model’s own typical activation size, not in raw units, so the same push strength means the same thing across all nine models. The push itself is

$$h' = h + \alpha \frac{\bar{h}}{\|\bar{h}\|} \hat{v},$$

where h is the story’s original activation, $\hat{v}$ is the unit-normalized target direction, $\|\bar{h}\|$ is that model’s mean residual-stream norm from Tab. 1, and α is the strength we choose.

At the concept level, the model answers a direct yes-or-no question: is this protagonist high in warmth (or competence)? Five strengths are tested, $\alpha \in \{-0.10, -0.05, 0, +0.05, +0.10\}$, spanning a strong pull toward “low” to a strong pull toward “high.” This range was fixed after an initial sweep up to ± 0.50 on Gemma-3-12B and Gemma-3-27B saturated the model’s response instead of shifting it cleanly, the same kind of breakdown at high steering strength that Turner et al. (2023) report for activation steering more generally; the narrower range was then applied uniformly to all nine models. For each strength, the direction is applied to 10 held-out topics it was never built from, separate from the 40 topics used to build it. For the Gemma-4 and Qwen3.6 checkpoints this yes-or-no question is presented through each model’s native chat template; the other four models receive it as raw text, matching how each checkpoint was run throughout.

We record how much more the model favors Yes over No at each strength, the next-token logit margin:

$$\Delta = \text{logit}(\text{Yes}) - \text{logit}(\text{No}),$$

where $\text{logit}(\text{Yes})$ and $\text{logit}(\text{No})$ are the model’s raw output scores for the two answer tokens before any softmax normalization, so a positive Δ means the model favors Yes and a larger Δ means a stronger preference. Throughout, the raw dense directions define the primary causal comparisons reported here; the PCA-denoised vectors serve as a robustness check (Supplementary, PCA Denoising).

A single steering result does not tell us whether the effect is specific to warmth or competence, or whether it

would show up for almost any direction we pushed along. To check this, we compare six different directions and see how each one shifts the model’s answer:

- **Dense target direction:** The same warmth or competence direction used everywhere else in this study, the difference between the average high- and low-condition activations. This is the main direction we are actually trying to validate.
- **GS2Decoded Gemma Scope 2 SAE direction:** Gemma Scope 2 (Google DeepMind Language Model Interpretability Team, 2025) is a separate tool from Google DeepMind that breaks a model’s internal activity down into thousands of smaller, more specific pieces. We compute the same warmth/competence contrast inside this alternative breakdown, then translate it back into the model’s own activation space, to see whether steering still works when built from a completely different representation.
- **GS2 Axis-specific component:** Within that breakdown, some pieces move together for both warmth and competence, and some move only for whichever concept is being tested. This direction keeps only the part specific to the one concept being tested.
- **GS2 Component shared between the two axes:** The opposite of the direction above, only the part that moves for both warmth and competence together, the piece that is not specific to either concept on its own.
- **GS2 Opposing concept axis:** When testing warmth, we steer along the competence direction instead, and vice versa when testing competence, to see whether an unrelated concept’s direction moves the answer too.
- **Random orthogonal control:** A direction chosen at random and adjusted so it does not overlap at all with the real target direction, our chance-level baseline.

The dense target direction and a random orthogonal control were both tested on all nine models; the other four were tested only on Gemma-3-12B and Gemma-3-27B. Those four, marked **GS2** above, are built from a Gemma Scope 2 decomposition, and that tool covers only the Gemma-3 generation (Google DeepMind Language Model Interpretability Team, 2025), so they could not be extended to Gemma-4, Qwen, or Llama. The nine-model random control is a differently constructed, calibrated set of many random directions per model rather than the single orthogonalized vector compared here, but it serves the same chance-level baseline role. If steering only shifts the answer along the real target direction, and not along the unrelated ones, that supports the effect being specific to warmth or competence rather than an artifact of pushing the residual stream in any direction. Uncertainty throughout is estimated by paired topic bootstrap resampling.




Figure 2: Residual-stream concept geometry (schematic). Low-warmth and high-warmth stories occupy different regions of this space, and the mean-difference direction $v = \bar{h}_{\text{high}} - \bar{h}_{\text{low}}$ (blue arrow) is what separates them. This panel shows only the warmth axis, so its v is the vector written v_W in Methods, with competence v_C defined the same way. Steering acts like grabbing the wheel of a car that is already driving down one street and turning it until the car crosses onto the next street over: the red dashed arrow, $h' = h + \alpha \hat{v}$, shows this same kind of forced redirection, pushing a story’s representation from the low-warmth region, across the decision boundary, into the high-warmth region. Concretely, this is done by adding a scaled copy of that same high-warmth direction to the story’s current representation. Illustrative synthetic activations, not measured data.

Steering Hiring Decisions. The concept-level test shows the model will move its stated judgment when pushed; whether that carries over to an actual hiring decision is a separate question. To answer it, we bring the callback task into the same steering setup. The hiring prompt pairs each of the 282 applicant first names from Gallo et al. (2024) with one fixed administrative-assistant application, identical in education, experience, and skills across every trial, so the name is the only signal that changes. The model reads the application and answers a single question, “would you recommend calling this candidate back for an interview? Answer with a single word: Yes or No.” We score that answer with the callback margin,

$$\Delta_{\text{callback}} = \text{logit}(\text{Yes}) - \text{logit}(\text{No}),$$

the same Yes-versus-No logit difference used at the concept level, now read from the hiring prompt rather than the concept-judgment prompt. Demographic labels stay out of the prompt. Race and gender are never shown to the model; each name is tagged afterward from the Gallo et al. (2024) coding, mirroring how human correspondence studies read group membership from names alone (Bertrand and Mullainathan, 2004). We then apply the same steering hooks used at the concept level while the model works through these applications, over a

60-name subset of the roster. Strength is swept broadly at $\alpha \in \{\pm 0.25, \pm 0.50\}$ on all nine models and refined at $\{\pm 0.05, \pm 0.10\}$ on every checkpoint except Llama-3.1-8B and Qwen3-14B, so that a push on the warmth or competence direction can be traced directly to a shift in the callback margin.

Names, Human Ratings, and Disparity. Independently of the hiring prompt, each applicant name is embedded in a neutral carrier sentence, and its residual-stream activation is projected onto the warmth and competence directions to yield a per-name probe score. We compare these scores against the aggregated human ratings behind them (24,220 rater judgments from 787 raters across ten source studies) using Spearman rank correlation, which tells us whether the model’s internal sense of a name’s warmth or competence lines up with how people actually perceive that name. Group-level callback disparities are then computed as differences in mean callback margin between name groups (race: 47 Black-signaling versus 180 White-signaling names; gender: 154 female versus 115 male), standardized by the within-model standard deviation of the callback margin so that models with different logit scales can be compared on equal footing. Human reference gaps come from the published per-name callback rates, aggregated over the 149 names that match the benchmark on first name. Finally, we test whether the internal warmth and competence representations actually carry the name-group signal into the callback decision, using bootstrap mediation along the path name group $\rightarrow$ probe score $\rightarrow$ callback margin (5,000 resamples, seed 20260527, $n = 269$ names with complete data, race subset $n = 227$), reporting indirect effects with 95% confidence intervals for all sixteen model-by-grouping-by-probe combinations.

Results

Discussion and Conclusion

That large language models trained on human-generated text would encode warmth and competence as internally recoverable representations is, in one sense, not surprising: these models distil patterns from the same texts in which warmth and competence shape social perception. What the present results establish empirically is that this encoding is linear, geometrically stable, causally functional at the concept level, and shared across architecturally distinct models from independent laboratories. Each of these properties matters: linearity makes the representation accessible to lightweight probes and steering interventions; causal functionality means the representation participates in explicit judgements about warmth and competence; and




Figure 3: Warmth and competence directions share model-dependent geometry. Each panel projects the same 200 synthetic stories onto raw dense probe-layer directions in an oblique basis. The arrow angle equals $\arccos[\cos(w_W, v_C)]$, so smaller angles indicate stronger geometric alignment. All panels use model-internal z-scores, shared axis limits, and equal x/y scaling.

cross-architecture convergence suggests the finding reflects a property of training data and the objectives of language modelling rather than a quirk of any single design.

The scale comparison between Gemma-3-12B and Gemma-3-27B complicates this picture in an instructive way. Both models encode warmth and competence with comparable precision, and the 27B model’s probe scores correlate slightly more strongly with human warmth and competence ratings. Yet the causal chain from concept representation to hiring decision differs sharply between them. At 12B, increasing the warmth representation produces a large, approximately linear increase in callback inclination (up to +2.37 logit units at $\alpha = +0.10$). At 27B, the equivalent intervention produces a modest positive shift at low strength (+1.97 at $\alpha = +0.05$) that collapses to a large negative shift at slightly higher strength (−2.66 at $\alpha = +0.10$). This non-monotone response is not saturation: the 27B baseline already sits at $P(\text{Yes}) = 0.76$, leaving room for upward movement, and the response is a sign reversal rather than a plateau.

A more plausible account is that the 27B decision layer is organised by multiple interacting representations; the concept direction, when amplified beyond a narrow window, begins to interfere with those interactions rather than reinforcing them.

A related puzzle emerges from the cross-model steerability analysis. Among the models with local-regime mediation estimates, Gemma-3-12B has the largest normalized warmth steerability (0.236), yet it shows no significant hiring-level mediation path. Llama-3.1-8B has the smallest value within that subset (0.029) and shows the strongest hiring-level indirect effect (bootstrap IE = +0.190). This steerability paradox illustrates that the capacity to shift a model’s representations via linear probing does not imply that those representations causally govern the downstream decision. Models in which representations are most cleanly isolable may be precisely those in which decision processes operate through parallel mechanisms that bypass the steering target.

The warmth anti-alignment observed in Llama-3.1-

Model	Probe layer (frac)	ρ (warmth, human)	ρ (competence, human)	ρ (callback, warmth)	ρ (callback, competence)
Gemma-3-12B	31 (0.65)	+0.366***	+0.239***	+0.150*	+0.149*
Gemma-3-27B	40 (0.65)	+0.396***	+0.272***	-0.160**	-0.156**
Llama-3.1-8B	20 (0.62)	-0.300***	-0.062 n.s.	+0.092 n.s.	+0.053 n.s.
Gemma-4-12B	31 (0.65)	+0.020 n.s.	+0.222***	-0.110 n.s.	-0.124*
Gemma-4-26B-A4B	19 (0.63)	-0.204***	-0.044 n.s.	+0.510***	+0.373***
Gemma-4-31B	39 (0.65)	+0.267***	+0.293***	+0.024 n.s.	+0.337***
Qwen3-14B	26 (0.65)	-0.193**	+0.465***	+0.195**	+0.426***
Qwen3.6-27B	42 (0.66)	+0.186**	+0.250***	+0.302***	+0.220***
Qwen3.6-35B-A3B	26 (0.65)	+0.211***	+0.131*	+0.344***	+0.197***

Table 2: Name-level correlation between probed warmth/competence and human ratings. For each of the 282 rated applicant names, the model’s residual-stream activation at the probe layer is projected onto the warmth and competence direction vectors. “Probe layer” gives the selected layer index and its resolved fraction of total depth (targeted probe_layer_frac = 0.66; the reported fraction is the integer-rounded layer divided by model depth). Warmth/competence columns correlate the probe projection with human warmth/competence ratings; callback columns correlate the probe projection with the model’s own unsteered callback margin. All correlations are Spearman ρ . * $p < .05$, ** $p < .01$, *** $p < .001$, n.s. not significant.

8B and Qwen3-14B adds a further layer of complexity. Despite high cross-model story agreement (per-story Spearman $\rho \approx 0.76$ – 0.77 between Gemma and each of the other models), the direction that Llama and Qwen organise as “high warmth” in activation space is negatively correlated with human warmth ratings ($\rho = -0.300$ and -0.193 respectively). The models converge on the same relative ordering of stories yet that shared ordering is inverted relative to the human scale. We interpret this as evidence that instruction tuning has restructured the warmth axis in these architectures: reinforcement learning from human feedback may have adjusted representations associated with stereotypically warm social groups in a direction that happens to invert the relationship to the original pretraining warmth axis.

The demographic disparity results at Gemma-3-27B illustrate the practical consequence of this representational reorganisation. The model strongly advantages Black-signalling names relative to White-signalling names (+1.18 SD), directly reversing the classical correspondence-study finding. The gender gap direction, by contrast, reproduces the human pattern. The asymmetry between race and gender suggests that instruction tuning has applied differential corrections to different demographic axes, producing a model that is not more equitable overall but differently inequitable: the racial hierarchy is reversed rather than eliminated, and at substantially larger magnitude than the original human disparity. This finding points to a fundamental limitation of outcome-level bias mitigation: correcting behavioral outputs without understanding the representational substrate can shift rather than remove the underlying imbalance.

Together, the results suggest that LLMs have internalised a structure of social perception broadly analogous to the SCM, emerging from training on human-generated text rather than through any explicit design choice. The connection from that internal structure to

actual hiring decisions is not, however, simple or consistent across models. This finding points to a methodological implication for bias auditing: probing internal representations and measuring output disparities are complementary but not interchangeable, and a model that performs well on one measure may not perform well on the other. It also opens a path to addressing the problem at the source by intervening on the representation rather than filtering outputs after the fact, and provides a methodological template for connecting the computational tools of mechanistic interpretability to the empirical tradition of correspondence studies in the social sciences.

Limitations. This study has several limitations, listed individually here so each concern can be weighed against its own evidence rather than folded into continuous prose:

- **Circularity in stimulus generation:** The concept story corpus was generated by a large language model, raising a concern about circularity in a study that probes language model representations. The generating model (Claude, Anthropic) and the probed models (Gemma, Qwen, Llama) come from different organizations and architectures, so the warmth/competence contrasts are not trivially reproduced by a shared generator. Even so, AI-generated stimuli could share stylistic properties that inflate probe accuracy across model families; replication with human-authored or template-based stimuli is desirable.
- **Limited topic coverage:** The corpus covers 50 topics. Broader topic coverage and larger per-condition samples would further test how far the directions generalize.
- **Moderate probe-to-human alignment:** The probe-to-human validation correlations are moderate (warmth $\rho = +0.366$ and $+0.396$ for 12B and 27B;

competence $\rho = +0.239$ and $+0.272$), indicating that the model’s internal probes capture a component of human social perception rather than replicating it faithfully.

- **Inverted warmth construct in two models:** The warmth anti-alignment in Llama-3.1-8B ($\rho = -0.300$) and Qwen3-14B ($\rho = -0.193$) means that for these models the mean-difference direction captures a construct systematically inverted relative to the human warmth scale; results for these models should be interpreted with this inversion in mind.
- **A fixed probe-layer heuristic, not validated per model:** We set `probe_layer_frac` = 0.66 for every model based on the relative depth reported for emotion representations in Sofroniew et al. (2026), applying it as a static, prespecified choice rather than testing separability at that depth on each model individually. A full-depth layer sweep on the training story corpus shows that this heuristic layer rarely coincides with the layer of peak separability: at Gemma-3-12B the selected layer reaches a warmth Cohen’s d of 2.68, while the best-separated layer in the same sweep reaches 6.27; Gemma-3-27B (2.95 versus 5.10) and Gemma-4-31B (7.56 versus 11.49) show comparable gaps. This mismatch does not fully explain the weak or inverted name-level correlations reported above: Gemma-4-12B already reaches a high separability score at its selected layer ($d = 8.46$), yet the resulting warmth probe fails to correlate with human name ratings ($\rho = +0.020$, n.s.). Story-level separability, which guided layer selection, therefore does not guarantee that the resulting direction generalizes to bare applicant names, a case the training stories never present. Validating or tuning the probe layer per model, rather than transferring a single depth from a different study, is likely more principled; whether averaging or weighting probe projections across several separable layers would recover a stronger name-level signal than the single fixed layer is a question we leave open.
- **Ceiling effects in cross-validated accuracy:** 5-fold and topic-holdout accuracy reach 1.000 with zero variance across folds for every model and both axes. This is consistent with the large Cohen’s d values reported above (for example $d = 2.68$ for warmth at Gemma-3-12B), but a ceiling this uniform is also expected on statistical grounds alone: with only 100 stories per axis and residual-stream vectors of several thousand dimensions, near-perfect linear separability is likely regardless of whether the direction captures a meaningful concept. Topic-holdout accuracy was designed specifically to close the topic-vocabulary leakage route available to plain 5-fold splitting, described above, yet it saturates at the same ceiling; each topic-holdout fold still contains only around ten

held-out topics against several-thousand-dimensional vectors, so the same high-dimension, low-sample-size argument applies there as well, and the stricter split does not by itself rule it out. Cross-validated accuracy at ceiling therefore cannot distinguish between models or axes on its own, and should be read alongside the effect-size and split-half checks rather than as independent evidence of signal strength.

- **Fragile causal effect at scale:** The causal warmth steering effect at the hiring level is positive and strong at 12B but non-monotone and fragile at 27B, so the causal claim does not generalize across scale without qualification.
- **Callback-margin quantization:** Callback margins are quantized to a 0.125-unit grid due to bf16 precision in the model’s output logits; this quantization is inherent to bf16 inference and is not resolved by float32 subtraction. We attempted a partial fix by casting the Yes/No logits to float32 before subtracting them, which removes rounding in the subtraction step itself, but the two operands are already bf16-quantized by the time they reach that subtraction, so the 0.125 grid persists regardless. A complete fix would require running inference itself in float32, which roughly doubles GPU memory usage and is likely infeasible for the 27B checkpoint on the available cluster hardware; we did not attempt it.
- **Narrow callback variance at some checkpoints:** At Gemma-3-12B the callback margin standard deviation is 0.14 with only seven unique values, making group-level disparity estimates at 12B unreliable; the 12B R4 results should not be cited as empirical findings. Llama-3.1-8B shows a comparably narrow spread ($SD = 0.12$, 12 unique values), so its group-level gaps should be read with the same caution. At Gemma-3-27B the larger variance ($SD = 0.41$, 18 unique values) makes group gaps of approximately 0.5 SD detectable, but the grid structure introduces discrete steps in all reported gap magnitudes. Qwen3-14B ($SD = 0.35$, 17 unique values) falls in the same usable range as 27B.
- **Uncorrected mediation tests:** The sixteen mediation tests are presented uncorrected for multiple comparisons; only the Llama race-warmth indirect effect survives a Bonferroni-corrected threshold, and the smaller significant paths should be treated as suggestive rather than confirmed.
- **Simplified hiring prompt:** The hiring prompt is simplified relative to real-world screening contexts (single role, fixed qualifications), and results may differ under richer evaluation settings or different job types.
- **Open architectural question:** The architectural source of the elevated warmth/competence vector cosine observed in Gemma-3 relative to Qwen3 and

Llama-3.1 across all network depths remains an open question.

- **Steering-strength range calibrated on two models:** We fixed the local steering range after testing saturation on Gemma-3-12B and Gemma-3-27B only, then applied it uniformly to all nine models; a more thorough methodology would calibrate this range separately for each model rather than reusing one pair’s calibration across architectures that differ substantially in scale.

Future Work. The immediate next step is expanding the hiring stimuli to additional job roles to test whether results generalise beyond administrative assistant. The SAE decomposition performed here on Gemma-3 can be extended to Llama-3.1 using the Llama Scope sparse autoencoders (He et al., 2024), enabling a model-family comparison at the feature level. Validating the concept probes against a human-authored warmth/competence dataset would address the circularity concern associated with AI-generated stimuli. On the social-science side, extending the hiring stimuli beyond names to category-level signals (age, disability, religion) would broaden the scope of the benchmark comparison and test whether the representational account generalises beyond race and gender. Translating these findings into practical debiasing interventions and evaluating their robustness to adversarial prompting is a natural next step for applied work.

Warmth and competence representations emerge with depth
(nine open-weights models grouped by model family and generation)




Figure 4: Warmth and competence directions emerge robustly across depth. Layer-wise Cohen's d across nine open-weights models. The dotted vertical line marks the fixed probe-layer fraction of 0.66.




Figure 5: Normalized concept-level steerability varies across nine checkpoints. Each curve shows the change in the held-out concept margin Δ (the Yes-versus-No logit difference defined in Methods) produced by target-direction steering, divided by the baseline high-versus-low gap in Δ for the same model and axis. Parentheses in the legend report the maximum positive steering coefficient, $\alpha = +0.10$, expressed relative to the mean residual norm. Normalization reduces raw logit-scale differences but does not establish direction specificity; calibrated random and cross-axis controls are interpreted in the text.




Figure 6: Activation steering shifts hiring callback recommendations. Callback-margin responses to warmth and competence steering in the hiring evaluation setting.




Figure 7: Matched examples show that positive concept steering can move callback decisions in either direction. Rows show warmth and competence; columns compare Gemma-3-12B, whose mean decision moves from No to Yes, with Gemma-3-27B, whose mean decision moves from Yes to No. All panels use the same 60 applications and $\alpha = +0.10$. Gray dashed paths show baseline outcomes, blue paths show steered outcomes, and bottom annotations summarize individual-name transitions. The panels were selected to display both transition directions under matched conditions, not to estimate their prevalence. Endpoint margins, strengths, and transition counts are empirical; connector geometry is schematic. Complete results for all nine models and both axes appear in Tab. S.1.




Figure 8: Dense steering dose-response across models. Concept-level dose-response curves for warmth and competence directions.




Figure 9: Model callback disparities relative to the human benchmark. Race and gender callback-margin gaps for the model evaluations compared with the correspondence-study benchmark.




Figure 10: Bootstrap mediation of name-group callback disparities through warmth and competence probe scores. Indirect effects with 95% confidence intervals for the pathway name group → probe score → callback margin; filled markers indicate intervals that exclude zero.

References

- An, Haozhe et al. (2024). *Do Large Language Models Discriminate in Hiring Decisions on the Basis of Race, Ethnicity, and Gender?* DOI: 10.48550/arXiv.2406.10486. arXiv: 2406.10486 [cs.CL].
- An, Jiafu et al. (2025). “Measuring gender and racial biases in large language models: Intersectional evidence from automated resume evaluation”. In: *PNAS Nexus* 4.3, pgaf089. ISSN: 2752-6542. DOI: 10.1093/pnasnexus/pgaf089. URL: <https://doi.org/10.1093/pnasnexus/pgaf089>.
- Bertrand, Marianne and Sendhil Mullainathan (2004). “Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination”. In: *American economic review* 94.4, pp. 991–1013.
- Cuddy, Amy J. C., Susan T. Fiske, and Peter Glick (2008). “Warmth and Competence as Universal Dimensions of Social Perception: The Stereotype Content Model and the BIAS Map”. In: *Advances in Experimental Social Psychology*. Vol. 40. Academic Press, pp. 61–149. DOI: 10.1016/S0065-2601(07)00002-0.
- Elhage, Nelson et al. (2021). *A Mathematical Framework for Transformer Circuits*. Transformer Circuits Thread. URL: <https://transformer-circuits.pub/2021/framework/index.html>.
- Fiske, Susan T., Amy J. C. Cuddy, and Peter Glick (2007). “Universal Dimensions of Social Cognition: Warmth and Competence”. In: *Trends in Cognitive Sciences* 11.2, pp. 77–83. DOI: 10.1016/j.tics.2006.11.005.
- Fiske, Susan T. et al. (2002). “A Model of (Often Mixed) Stereotype Content: Competence and Warmth Respectively Follow from Perceived Status and Competition”. In: *Journal of Personality and Social Psychology* 82.6, pp. 878–902. DOI: 10.1037/0022-3514.82.6.878.
- Gaebler, Johann D. et al. (2024). “Auditing Large Language Models for Race and Gender Disparities: Implications for AI-Based Hiring”. In: *Socius* 11. DOI: 10.1177/23794607251320229.
- Gallo, Marcos et al. (2024). “Perceived Warmth and Competence Predict Callback Rates in Meta-Analyzed North American Labor Market Experiments”. In: *PLOS ONE* 19.7, e0304723. DOI: 10.1371/journal.pone.0304723.
- Gemma Team (2025). *Gemma 3 Technical Report*. arXiv: 2503.19786 [cs.CL]. URL: <https://arxiv.org/abs/2503.19786>.
- Google DeepMind Language Model Interpretability Team (2025). *Gemma Scope 2: Helping the AI Safety Community Deepen Understanding of Complex Language Model Behavior*. Google DeepMind Blog.
- Grattafiori, Aaron et al. (2024). *The Llama 3 Herd of Models*. arXiv: 2407.21783 [cs.AI]. URL: <https://arxiv.org/abs/2407.21783>.
- He, Zhengfu et al. (2024). *Llama Scope: Extracting Millions of Features from Llama-3.1-8B with Sparse Autoencoders*. arXiv: 2410.20526 [cs.CL]. URL: <https://arxiv.org/abs/2410.20526>.
- Mikolov, Tomas et al. (2013). *Efficient Estimation of Word Representations in Vector Space*. arXiv: 1301.3781 [cs.CL].
- Nanda, Neel and Joseph Bloom (2022). *TransformerLens*. <https://github.com/TransformerLensOrg/TransformerLens>.
- Olah, Chris et al. (2020). “Zoom In: An Introduction to Circuits”. In: *Distill*. DOI: 10.23915/distill.00024.001.
- Park, Kiho, Yo Joong Choe, and Victor Veitch (2024). *Linear Representation Hypothesis and the Geometry of Large Language Models*. arXiv: 2311.03658 [cs.LG].
- Sofroniew, Nicholas et al. (2026). *Emotion Concepts and Their Function in a Large Language Model*. DOI: 10.48550/arXiv.2604.07729. arXiv: 2604.07729 [cs.CL].
- Sun, Mingjie et al. (2024). “Massive Activations in Large Language Models”. In: *arXiv preprint arXiv:2402.17762*. URL: <https://arxiv.org/abs/2402.17762>.
- Turner, Alexander Matt et al. (2023). *Activation Addition: Steering Language Models Without Optimization*. arXiv: 2308.10248 [cs.LG].
- Vaswani, Ashish et al. (2017). “Attention Is All You Need”. In: *Advances in Neural Information Processing Systems* 30.
- Xiao, Guangxuan et al. (2024). “Efficient Streaming Language Models with Attention Sinks”. In: *International Conference on Learning Representations (ICLR)*. arXiv:2309.17453. URL: <https://arxiv.org/abs/2309.17453>.
- Yang, An et al. (2025). *Qwen3 Technical Report*. arXiv: 2505.09388 [cs.CL]. URL: <https://arxiv.org/abs/2505.09388>.
- Zou, Andy et al. (2023). *Representation Engineering: A Top-Down Approach to AI Transparency*. arXiv: 2310.01405 [cs.LG].

Supplementary Materials

Details on Methods

Compute and Hardware

All experiments ran on two compute environments. The SCKKN cluster (Universität Konstanz) supplied NVIDIA RTX PRO 6000 Blackwell GPUs (96 GB VRAM), which ran the Gemma-3, Llama-3.1, and Qwen3 jobs (Gemma-3-12B, Gemma-3-27B, Llama-3.1-8B, Qwen3-14B) and most of the extraction and steering jobs for the Gemma-4 and Qwen3.6 checkpoints. Gemma-4-12B is a partial exception: its Stage 1 activations were first extracted on an NVIDIA L40, and a later retry ran on an L40S instead. A dedicated reproducibility audit on the exact L40 hardware class found that both device classes agree on which layer peaks for warmth and competence (layers 26 and 27, respectively), but differ non-negligibly at some layers (maximum competence Cohen’s d difference of 0.914, cosine difference of 0.095). The exact-L40 run, hardware-matched to the original Stage 1 extraction, is the one reported throughout the main text. Select calibrated-steering replications for the Gemma-4 and Qwen3.6 checkpoints (all three Gemma-4 models and both Qwen3.6 models) instead ran on NVIDIA H100 GPUs at the Collective Computational Unit (CCU). The authors acknowledge support by the local computing resources through the core facility SCKKN. The authors additionally acknowledge the Collective Computational Unit (CCU) for providing H100 GPU access used in part of this work.

PCA Denoising

The warmth and competence direction vectors share a substantial common component that reflects general narrative valence: stories that portray a protagonist positively (whether warmly or competently) occupy a similar region of the residual stream, and the same holds for negatively valenced stories. To separate conceptual content from this shared evaluative polarity, we follow the neutral-corpus denoising procedure of Sofroniew et al. (2026). We build a neutral corpus of 1,500 Wikipedia introduction paragraphs, length-matched to the concept stories (90–200 words) and filtered to exclude emotionally charged or violent content, and extract residual-stream activations at the same layer and pooling rule used for the concept stories. We then fit principal component analysis (PCA) on these neutral activations and project out, from both direction vectors, the top components that together explain at least 50 percent of neutral variance. For Gemma-3-12B, the first principal component alone explains 56.1 percent of neutral variance, and removing it lowers the cosine between the warmth and competence vectors from 0.749 to 0.530. For Gemma-3-27B, 43 components are needed to cross the 50 percent threshold (50.2 percent explained), lowering the inter-axis cosine from 0.708 to 0.487. The remaining post-denoising cosine, 0.53 and 0.49 respectively, is close to the genuine warmth-competence correlation observed in human ratings ($r \approx 0.61$; Fiske et al. 2002), which suggests that the residual overlap reflects real conceptual co-variation rather than a failure of the denoising step. The denoised vectors serve as a robustness check; the raw dense directions define the primary cross-model causal comparisons and the callback-transition census reported in the main text.

Additional Results

Complete Positive-Steering Callback Transitions

Tab. S.1 reports every model-axis condition used in the positive-endpoint comparison. The table is the complete result; the matched panels in Fig. 7 are illustrative examples selected under a common family, application set, and steering strength. Gemma-3-12B, Gemma-3-27B, and the Qwen3.6 and Gemma-4 checkpoints use $\alpha = +0.10$. Llama-3.1-8B and Qwen3-14B use their available broad-regime endpoint of $\alpha = +0.50$, so raw effect sizes should not be compared directly across all rows.

Positive steering crosses the mean decision boundary in both directions among the Gemma-3 models and from No to Yes for Llama-3.1-8B. The other six models remain in the Yes category for every name, even when the mean margin decreases. This pattern reflects two distinct mechanisms. First, a concept direction constructed from high-minus-low stories is not constrained to align positively with the callback readout. Second, models beginning far from zero can move substantially without changing the categorical recommendation. The Gemma-3-27B response adds a

third complication because its local dose-response reverses between +0.05 and +0.10, indicating a narrow nonlinear intervention regime.

Model	α	Baseline	Warmth			Competence		
			Endpoint	Δ margin	Transitions	Endpoint	Δ margin	Transitions
Gemma-3-12B	+0.10	-0.194 (N)	+2.175 (Y)	+2.369	54 N→Y; 6 T→Y	+0.037 (Y)	+0.231	19 N→N, 13 N→T, 22 N→Y
Gemma-3-27B	+0.10	+1.169 (Y)	-1.490 (N)	-2.658	60 Y→N	-1.527 (N)	-2.696	1 T→N, 2 T→T, 3 T→Y
Llama-3.1-8B	+0.50	-2.098 (N)	+1.073 (Y)	+3.171	60 N→Y	+0.072 (Y)	+2.170	60 Y→N
Gemma-4-12B	+0.10	+17.188 (Y)	+18.237 (Y)	+1.049	60 Y→Y	+18.715 (Y)	+1.527	15 N→N; 10 N→T; 35 N→Y
Gemma-4-26B-A4B	+0.10	+21.245 (Y)	+21.307 (Y)	+0.062	60 Y→Y	+20.865 (Y)	-0.380	60 Y→Y
Gemma-4-31B	+0.10	+25.615 (Y)	+25.173 (Y)	-0.442	60 Y→Y	+24.801 (Y)	-0.814	60 Y→Y
Qwen3-14B	+0.50	+1.704 (Y)	+2.304 (Y)	+0.600	60 Y→Y	+1.321 (Y)	-0.383	60 Y→Y
Qwen3.6-27B	+0.10	+5.346 (Y)	+6.542 (Y)	+1.196	60 Y→Y	+5.879 (Y)	+0.533	60 Y→Y
Qwen3.6-35B-A3B	+0.10	+3.031 (Y)	+3.998 (Y)	+0.967	60 Y→Y	+3.465 (Y)	+0.433	60 Y→Y

Table S.1: Complete positive-endpoint callback transitions. Mean baseline and steered callback margins are followed by their decision category in parentheses. Transition cells list every nonzero individual-name transition among No (N), Tie (T), and Yes (Y); each cell sums to 60.

Name-Level Warmth/Competence by Marginal Group

Tab. S.2 breaks the same 282-name probe scores down by race (Black/White) and gender (Female/Male) as separate marginal axes, following Gallo et al. (2024)’s own grouping convention. Model warmth and competence values are raw, unnormalized projections onto the concept direction, so they can only be compared within a model, not across the nine rows.

Name-Level Warmth/Competence by Crossed Race $\times$ Gender

Tab. S.3 instead crosses race and gender into four joint cells (Black-Female, Black-Male, White-Female, White-Male), recovering the interaction structure that the marginal breakdown above collapses. Group sizes are small and uneven, from 9 names in each Black cell to 84 in White-Female, so the crossed cell means should be read as descriptive rather than as adequately powered group comparisons.

Bootstrap Mediation, All Nine Models

The main text reports bootstrap mediation for four models (Gemma-3-12B, Gemma-3-27B, Llama-3.1-8B, Qwen3-14B) and builds the steerability-paradox argument on those sixteen tests. Tab. S.4 extends the identical procedure, same bootstrap resamples, same seed, same standardization, to the remaining five models (Gemma-4-12B, Gemma-4-26B-A4B, Gemma-4-31B, Qwen3.6-27B, Qwen3.6-35B-A3B), for a combined thirty-six tests across all nine models. Fourteen of the thirty-six indirect effects exclude zero at the uncorrected 95% level; under a Bonferroni threshold across all thirty-six ($\alpha = 0.05/36$), only the Llama-3.1-8B race–warmth path from the main text survives, so every other entry in this table, old or new, remains suggestive rather than confirmed.

Beyond those four models, the remaining five complicate the paradox in one specific way. Both Gemma-3 models showed no significant mediation path on either axis, which the main text reads as evidence that Gemma’s cleanly steerable representations bypass the hiring decision rather than driving it. Gemma-4 breaks that pattern: all three Gemma-4 checkpoints show a significant race–competence indirect effect, and Gemma-4-26B-A4B additionally shows a significant gender–warmth path. Qwen3.6 shows a comparable spread of significant competence and warmth paths across its two checkpoints, consistent with the mediation profile already observed in Qwen3-14B. Read together, these results suggest that the absence of hiring-level mediation may be a property of the Gemma-3 generation specifically rather than of architectures with high concept-steerability in general; the four-model steerability paradox in the main text should be understood as scoped to that subset rather than as a claim about steerable models at large.

Story Corpus by Topic

Model	Group	<i>n</i>	Human callback	Model margin	Model warmth / competence
Gemma-3-12B	Black	47	0.183	-0.213	37605.34 / 40319.84
	White	180	0.171	-0.199	38362.85 / 41133.78
	Female	154	0.145	-0.165	38071.05 / 40816.69
	Male	115	0.181	-0.258	37949.40 / 40692.73
Gemma-3-27B	Black	47	0.183	1.665	26143.95 / 25183.93
	White	180	0.171	1.121	26600.90 / 25607.80
	Female	154	0.145	1.118	26448.04 / 25462.46
	Male	115	0.181	1.312	26333.69 / 25349.97
Llama-3.1-8B	Black	47	0.183	-2.020	-1.27 / -0.16
	White	180	0.171	-2.067	-1.34 / -0.17
	Female	154	0.145	-2.054	-1.31 / -0.16
	Male	115	0.181	-2.107	-1.32 / -0.17
Gemma-4-12B	Black	47	0.183	17.194	12.22 / 18.46
	White	180	0.171	17.172	12.29 / 18.96
	Female	154	0.145	17.208	12.23 / 18.70
	Male	115	0.181	17.175	12.30 / 18.74
Gemma-4-26B-A4B	Black	47	0.183	21.509	1.08 / 5.59
	White	180	0.171	21.095	0.90 / 5.47
	Female	154	0.145	21.464	1.00 / 5.51
	Male	115	0.181	20.902	0.92 / 5.54
Gemma-4-31B	Black	47	0.183	25.808	45.31 / 78.15
	White	180	0.171	25.577	45.52 / 78.77
	Female	154	0.145	25.788	45.34 / 78.40
	Male	115	0.181	25.408	45.43 / 78.38
Qwen3-14B	Black	47	0.183	1.758	-8.46 / 17.42
	White	180	0.171	1.700	-9.14 / 17.75
	Female	154	0.145	1.851	-8.64 / 17.76
	Male	115	0.181	1.537	-8.99 / 17.48
Qwen3.6-27B	Black	47	0.183	5.402	5.22 / 8.10
	White	180	0.171	5.328	5.19 / 8.20
	Female	154	0.145	5.384	5.19 / 8.09
	Male	115	0.181	5.284	5.19 / 8.24
Qwen3.6-35B-A3B	Black	47	0.183	2.981	0.17 / 0.39
	White	180	0.171	3.039	0.18 / 0.40
	Female	154	0.145	3.143	0.18 / 0.39
	Male	115	0.181	2.873	0.17 / 0.39

Table S.2: Name-level warmth/competence and callback margin by marginal demographic group. Race (Black/White) and gender (Female/Male) are treated as separate marginal axes, following Gallo and Hausladen’s own grouping convention. Model warmth/competence are raw (unnormalized) projections onto the concept direction and are not comparable in magnitude across models; only within-model group differences are meaningful.

Model	Race $\times$ Gender	n	Human callback	Model margin	Model warmth / competence
Gemma-3-12B	Black-Female	9	0.065	-0.181	37660.74 / 40372.96
	Black-Male	9	0.059	-0.278	37842.95 / 40574.00
	White-Female	84	0.129	-0.171	38486.83 / 41263.31
	White-Male	47	0.179	-0.301	38273.64 / 41035.44
Gemma-3-27B	Black-Female	9	0.065	1.514	26107.43 / 25164.89
	Black-Male	9	0.059	1.778	26274.15 / 25309.69
	White-Female	84	0.129	1.033	26642.11 / 25653.64
	White-Male	47	0.179	1.274	26446.95 / 25463.59
Llama-3.1-8B	Black-Female	9	0.065	-1.979	-1.26 / -0.15
	Black-Male	9	0.059	-2.069	-1.27 / -0.18
	White-Female	84	0.129	-2.067	-1.33 / -0.17
	White-Male	47	0.179	-2.104	-1.34 / -0.17
Gemma-4-12B	Black-Female	9	0.065	17.121	12.19 / 18.40
	Black-Male	9	0.059	17.266	12.22 / 18.51
	White-Female	84	0.129	17.173	12.26 / 18.96
	White-Male	47	0.179	17.108	12.36 / 19.00
Gemma-4-26B-A4B	Black-Female	9	0.065	21.840	1.14 / 5.56
	Black-Male	9	0.059	21.160	0.98 / 5.49
	White-Female	84	0.129	21.275	0.92 / 5.44
	White-Male	47	0.179	20.595	0.86 / 5.44
Gemma-4-31B	Black-Female	9	0.065	25.889	45.20 / 78.05
	Black-Male	9	0.059	25.781	45.41 / 78.30
	White-Female	84	0.129	25.817	45.45 / 78.78
	White-Male	47	0.179	25.037	45.58 / 78.76
Qwen3-14B	Black-Female	9	0.065	1.806	-8.39 / 17.61
	Black-Male	9	0.059	1.667	-8.57 / 17.11
	White-Female	84	0.129	1.839	-9.07 / 17.72
	White-Male	47	0.179	1.391	-9.43 / 17.45
Qwen3.6-27B	Black-Female	9	0.065	5.431	5.22 / 8.00
	Black-Male	9	0.059	5.361	5.30 / 8.20
	White-Female	84	0.129	5.368	5.17 / 8.10
	White-Male	47	0.179	5.181	5.20 / 8.28
Qwen3.6-35B-A3B	Black-Female	9	0.065	3.167	0.17 / 0.39
	Black-Male	9	0.059	2.903	0.17 / 0.39
	White-Female	84	0.129	3.174	0.18 / 0.40
	White-Male	47	0.179	2.809	0.18 / 0.40

Table S.3: Name-level warmth/competence and callback margin by crossed race $\times$ gender group. Applicant names are joined to Gallo et al. (2024)’s published race/gender labels by lowercase first name and matching study (`src/hiring_r4.py`). Model warmth/competence are raw projections and are not comparable in magnitude across models; only within-model comparisons across the four cells are meaningful.

Model	Grouping	Probe	a	b	IE	95% CI
Gemma-3-12B	Race (Black=1)	Warmth	-0.502	+0.171	-0.086	[-0.186, +0.019]
	Race (Black=1)	Competence	-0.501	+0.174	-0.087	[-0.188, +0.018]
	Gender (Female=1)	Warmth	+0.072	+0.236	+0.017	[-0.011, +0.056]
	Gender (Female=1)	Competence	+0.068	+0.240	+0.016	[-0.012, +0.055]
Gemma-3-27B	Race (Black=1)	Warmth	-0.391	-0.114	+0.045	[-0.008, +0.106]
	Race (Black=1)	Competence	-0.383	-0.108	+0.041	[-0.010, +0.102]
	Gender (Female=1)	Warmth	+0.093	-0.066	-0.006	[-0.027, +0.006]
	Gender (Female=1)	Competence	+0.096	-0.055	-0.005	[-0.025, +0.007]
Llama-3.1-8B	Race (Black=1)	Warmth	+0.519	+0.366	+0.190 [†]	[+0.111, +0.292]
	Race (Black=1)	Competence	+0.218	+0.185	+0.040 [†]	[+0.004, +0.090]
	Gender (Female=1)	Warmth	+0.027	-0.112	-0.003	[-0.022, +0.016]
	Gender (Female=1)	Competence	+0.208	-0.012	-0.003	[-0.035, +0.030]
Gemma-4-12B	Race (Black=1)	Warmth	-0.260	-0.109	+0.028	[-0.003, +0.069]
	Race (Black=1)	Competence	-0.515	+0.291	-0.150 [†]	[-0.295, -0.019]
	Gender (Female=1)	Warmth	-0.327	-0.119	+0.039	[-0.004, +0.079]
	Gender (Female=1)	Competence	-0.032	-0.024	+0.001	[-0.015, +0.011]
Gemma-4-26B-A4B	Race (Black=1)	Warmth	+0.476	+0.173	+0.082	[-0.038, +0.194]
	Race (Black=1)	Competence	+0.387	+0.339	+0.131 [†]	[+0.060, +0.231]
	Gender (Female=1)	Warmth	+0.213	+0.181	+0.038 [†]	[+0.007, +0.078]
	Gender (Female=1)	Competence	-0.073	+0.386	-0.028	[-0.076, +0.023]
Gemma-4-31B	Race (Black=1)	Warmth	-0.372	+0.143	-0.053	[-0.164, +0.005]
	Race (Black=1)	Competence	-0.465	+0.494	-0.230 [†]	[-0.483, -0.049]
	Gender (Female=1)	Warmth	-0.094	+0.088	-0.008	[-0.042, +0.002]
	Gender (Female=1)	Competence	+0.010	+0.186	+0.002	[-0.029, +0.026]
Qwen3-14B	Race (Black=1)	Warmth	+0.488	+0.165	+0.081 [†]	[+0.013, +0.156]
	Race (Black=1)	Competence	-0.255	+0.517	-0.132 [†]	[-0.217, -0.058]
	Gender (Female=1)	Warmth	+0.210	+0.014	+0.003	[-0.017, +0.024]
	Gender (Female=1)	Competence	+0.218	+0.347	+0.076 [†]	[+0.029, +0.141]
Qwen3.6-27B	Race (Black=1)	Warmth	+0.070	+0.411	+0.029	[-0.029, +0.106]
	Race (Black=1)	Competence	-0.169	+0.290	-0.049 [†]	[-0.103, -0.011]
	Gender (Female=1)	Warmth	-0.021	+0.415	-0.009	[-0.059, +0.047]
	Gender (Female=1)	Competence	-0.320	+0.383	-0.123 [†]	[-0.206, -0.061]
Qwen3.6-35B-A3B	Race (Black=1)	Warmth	-0.228	+0.454	-0.103 [†]	[-0.169, -0.046]
	Race (Black=1)	Competence	-0.097	+0.354	-0.034	[-0.080, +0.006]
	Gender (Female=1)	Warmth	+0.249	+0.257	+0.064 [†]	[+0.027, +0.113]
	Gender (Female=1)	Competence	+0.124	+0.234	+0.029 [†]	[+0.001, +0.067]

Table S.4: Bootstrap mediation of name-group callback disparities through warmth and competence probe scores, all nine models. Path: name group $\rightarrow$ probe score $\rightarrow$ callback margin, standardized coefficients, Baron–Kenny/Preacher–Hayes bootstrap (5,000 resamples, seed 20260527, $n = 269$, race subset $n = 227$). a is the coefficient of group on probe score; b is the coefficient of probe score on callback margin, controlling for group; IE = $a \times b$ is the indirect (mediated) effect. [†] marks a 95% bootstrap CI that excludes zero. 14 of 36 tests are significant at this uncorrected threshold; only the Llama-3.1-8B race–warmth path (main text) survives a Bonferroni correction across all 36 tests ($\alpha = 0.05/36$), so the remaining paths listed here, including all newer-model paths, should be read as suggestive rather than confirmed.

No.	Topic	Category	Words
1	A team meeting where a decision needs to be made under time pressure	Workplace	130–142
2	Receiving critical feedback from a supervisor in front of colleagues	Workplace	100–104
3	Presenting results that did not go as expected to senior management	Workplace	95–103
4	Discovering that a colleague made a serious mistake before the deadline	Workplace	99–103
5	Mentoring a junior colleague on their first week at the job	Workplace	97–105
6	Handling an escalating complaint from a dissatisfied customer	Workplace	93–100
7	Negotiating project scope with a client who keeps changing requirements	Workplace	91–92
8	Organising a company event with a limited budget and unclear instructions	Workplace	89–96
9	Taking over a project mid-way when the previous lead leaves unexpectedly	Workplace	90–101
10	Responding to an emergency system failure at 2am on call	Workplace	92–102
11	Being asked to do something you believe is wrong by a manager	Learning	91–102
12	Learning a completely new technical skill under pressure and alone	Learning	97–104
13	Debugging a complex problem with no documentation and no internet	Learning	94–99
14	Preparing for a high-stakes exam with only two days left	Learning	93–97
15	Trying to understand a subject that feels overwhelming at first	Learning	93–101
16	Explaining a technical concept to someone with no background in it	Learning	126–144
17	Helping a stranger who dropped their groceries in the street	Social	88–97
18	A difficult conversation with a family member about money	Social	90–100
19	Being the only person at a party who does not know anyone	Social	92–95
20	Comforting a friend who just received bad news	Social	89–102
21	Waiting for important medical test results	Health	92–109
22	Supporting a parent who is beginning to lose their independence	Health	90–95
23	Managing a chronic illness while trying to maintain a normal schedule	Health	88–96
24	Recovering from an injury when you were counting on being active	Health	90–105
25	Organising a neighbourhood clean-up with volunteers who cancel last minute	Community	129–140
26	Attending a community meeting where people strongly disagree	Community	91–100
27	Volunteering at a food bank on an unexpectedly chaotic shift	Community	89–99
28	Running for a small local position against an experienced opponent	Community	93–106
29	Advocating for a change at the local school as a parent	Community	90–101
30	Deciding whether to take a pay cut for a job you believe in	Finance	88–95
31	Realising mid-month that you have far less money than expected	Finance	90–104
32	Negotiating a salary for the first time in a new job offer	Finance	90–103
33	Dealing with a billing dispute that has lasted several months	Finance	89–103
34	Coaching a sports team through a long losing streak	Sport & Creative	129–143
35	Performing music in public for the first time and making a mistake	Sport & Creative	90–101
36	Finishing a creative project that has been stalled for months	Sport & Creative	97–107
37	Being a referee in a game where the players are friends	Sport & Creative	93–98
38	Entering a competition knowing you are not the favourite	Sport & Creative	91–107
39	Missing a connecting flight with no phone battery and no local currency	Travel	125–142
40	Navigating a city you have never been to, in a country whose language you do not speak	Travel	90–100
41	Being stranded overnight due to a transport strike with strangers	Travel	93–107
42	Starting over in a new country with no local contacts	Travel	91–99
43	Arriving at accommodation that is nothing like what was advertised	Travel	93–102
44	Realising your important data was lost and there is no backup	Technology	90–95
45	Trying to set up a complex technical system with no instructions	Technology	95–105
46	Being asked to fix someone else’s code under a tight deadline	Technology	91–100
47	Teaching an elderly relative how to use a new device they are afraid of	Technology	90–102
48	Apologising sincerely to someone you have genuinely hurt	Personal Milestones	94–104
49	Being the first person in your family to pursue a university degree	Personal Milestones	92–102
50	Caring for a newborn alone for the first time with no experience	Personal Milestones	90–98

Table S.6: Story corpus by topic ($n = 50$). Every topic contributes exactly one story to each of the four conditions (high warmth, low warmth, high competence, low competence), for 200 stories total, and every protagonist is nameless and referred to only as “they” throughout (see Story Preparation). The word-count range gives the shortest and longest of a topic’s four stories. The Category column groups topics into the ten domains named in Story Preparation; these labels are our own assignment for readability and are not a field in the underlying stimuli data.